


Hash Tables

Direct-address table



DIRECT-ADDRESS-SEARCH(T, k)

1 **return** $T[k]$


DIRECT-ADDRESS-INSERT(T, x)

1 $T[x.key] = x$


DIRECT-ADDRESS-DELETE(T, x)

1 $T[x.key] = \text{NIL}$

Hash table



Collision resolution by chaining



CHAINED-HASH-INSERT(T, x)

1 LIST-PREPEND($T[h(x.key)]$, x)

CHAINED-HASH-SEARCH(T, k)

1 **return** LIST-SEARCH($T[h(k)]$, k)

CHAINED-HASH-DELETE(T, x)

1 LIST-DELETE($T[h(x.key)]$, x)

Analysis of hashing with chaining:

- Given a hash table with m slots and n keys, define load factor $\alpha = n/m$: average number of keys per slot.
- For $j=0, 1, \dots, m-1$, denote the length of the list $T[j]$ by n_j , so that $E[n_j] = \frac{n}{m} = \alpha$.
- Assume each key is equally likely to be hashed into any slot: **independently uniform hashing (IUH)**.

- Thm 1: In a hash table in which collisions are resolved by chaining, an **unsuccessful** search takes expected time $\Theta(1 + \alpha)$ under IUH.

Proof:

- Under the assumption of IUH, any **un-stored key** is equally likely to hash to any of the **m** slots.
- The expected time to search unsuccessfully for a key k is the expected time to search to the end of list $T[h(k)]$, **which is exactly α** .
- Thus, the total time required is $\Theta(1 + \alpha)$. □

- Thm 2: In a hash table in which collisions are resolved by chaining, a successful search takes $\Theta(1 + \alpha)$ time on average under IUH.

Proof:

- Let the element being searched for equally likely to be any of the n elements stored in the table.
- The expected time to search successfully for a key:



- Elements before x in the list were inserted after x was inserted.
- We want to find the expected number of elements added to x 's list after x was added to the list.

- Let x_i be the **i**th element into the table, for $i = 1$ to n , and let $k_i = x_i.key$.
- Define $X_{ijq} = I\{ \text{search } x_i, h(k_i) = q, h(k_j) = q \}$.
- Under IUH, we have $\Pr\{ h(k_i) = q \} = \Pr\{ h(k_j) = q \} = 1/m$
 $\Pr\{\text{search } x_i\} = 1/n, \quad E[X_{ijq}] = 1/nm^2$.
- Let $Y_j = I\{x_j \text{ appear in a list prior the element being searched}\} = \sum_{q=0}^{m-1} \sum_{i=1}^{j-1} X_{ijq}$
- Let $Z = \sum_{j=1}^n Y_j$, the number of elements appear in the list prior to the being searched element.

$$\begin{aligned}
E[Z + 1] &= E\left[1 + \sum_{j=1}^n Y_j\right] \\
&= 1 + E\left[\sum_{j=1}^n \sum_{q=0}^{m-1} \sum_{i=1}^{j-1} X_{ijq}\right] \\
&= 1 + E\left[\sum_{q=0}^{m-1} \sum_{j=1}^n \sum_{i=1}^{j-1} X_{ijq}\right] \\
&= 1 + \sum_{q=0}^{m-1} \sum_{j=1}^n \sum_{i=1}^{j-1} E[X_{ijq}] \quad (\text{by linearity of expectation}) \\
\\
&= 1 + \sum_{q=0}^{m-1} \sum_{j=1}^n \sum_{i=1}^{j-1} \frac{1}{nm^2} \\
&= 1 + m \cdot \frac{n(n-1)}{2} \cdot \frac{1}{nm^2} \\
&= 1 + \frac{n-1}{2m} \\
&= 1 + \frac{n}{2m} - \frac{1}{2m} \\
&= 1 + \frac{\alpha}{2} - \frac{\alpha}{2n}.
\end{aligned}$$

Thus, the total time required for a successful search (including the time for computing the hash function) is:

$$\Theta\left(2 + \frac{\alpha}{2} - \frac{\alpha}{2n}\right) = \Theta(1 + \alpha).$$

The multiplication method:

$$0 < A < 1$$




Figure 11.4 The multiplication method of hashing. The w -bit representation of the key k is multiplied by the w -bit value $s = A \cdot 2^w$. The p highest-order bits of the lower w -bit half of the product form the desired hash value $h(k)$.

Knuth suggests that $A \approx (\sqrt{5} - 1) / 2 = 0.6180339887\dots$

Take $k = 123456$, $p = 14$, $m = 2^{14} = 16384$ and $w = 32$.

$$s = A \cdot 2^{32} = 2654435769$$

$$k \cdot s = 327706022297664 = (76300 \cdot 2^{32}) + 176112864$$

The 14 most significant bits of 176112864 is 67!

Random hashing

- With a fixed hash function, a malicious adversary can choose n keys that all hash to the same slot, yielding an average retrieval time of $\Theta(n)$.
- Any static hash function is vulnerable to such terrible worst-case behavior.
- The only effective way to improve the situation is to **choose the hash function randomly** in a way that is independent of the keys that are actually going to be stored.
- This approach is called **random hashing**. A special case of this approach, called **universal hashing**.

Universal Hashing

- $\mathcal{H} = \{ h: U \rightarrow \{0, \dots, m - 1\} \}$, which is a finite collection of hash functions.
- \mathcal{H} is called “*universal*” if for each pair of distinct keys $\ell, k \in U$, the number of hash functions $h \in \mathcal{H}$ for which $h(k) = h(\ell)$ is at most $|\mathcal{H}|/m$
- Define n_i = the length of list $T[i]$

Cor 3: Using universal hashing and collision resolution by chaining in an initially empty table with m slots, it takes $\Theta(s)$ expected time to handle any sequence of s **INSERT**, **SEARCH**, and **DELETE** operations containing $n = O(m)$ **INSERT** operations.

Proof:

- Each **INSERT** and **DELETE** op take constant time.
- $\alpha = \frac{n}{m} = O(1)$.
- The expected time for each **SEARCH** operation is $O(1)$.
- $\Pr_h\{h(k_1) = h(k_2)\} \leq 1/m$.
- By linearity of expectation, therefore, the expected time for the entire sequence of s operations is $O(s)$.

Achievable properties of random hashing

$\mathcal{H} = \{ h: U \rightarrow \{0, \dots, m - 1\} \}$, which is a finite collection of hash functions. h can be uniformly selected from \mathcal{H} .

- \mathcal{H} is “*uniform*” if for each key $k \in U$ and any slot q in the range $\{0, 1, \dots, m - 1\}$, $\Pr\{h(k) = q\} = 1/m$.
- \mathcal{H} is “*universal*” if for each pair of distinct keys $\ell, k \in U$,

$$\Pr_{h \in \mathcal{H}} \{h(k) = h(\ell)\} \leq \frac{1}{m}.$$

- \mathcal{H} is “ ϵ – *universal*” if for each pair of distinct keys $\ell, k \in U$,
$$\Pr_{h \in \mathcal{H}} \{h(k) = h(\ell)\} \leq \epsilon$$
- \mathcal{H} is “ d – *independent*” if for any d distinct keys $k_1, \dots, k_d \in U$ and any d slots q_1, \dots, q_d , $\Pr_h \{h(k_i) = q_i, i = 1 \dots d\} = 1/m^d$.

Design a universal class of hash functions:

- Let p be a prime number.

$$Z_p = \{0, \dots, p - 1\}, \quad Z_p^* = Z_p / \{0\}$$

- For any $a \in Z_p^*$ and $b \in Z_p$, define:

$$h_{a,b}(k) = ((ak + b) \bmod p) \bmod m$$

$$\mathcal{H}_{p,m} = \{h_{a,b} : a \in Z_p^* \text{ and } b \in Z_p\}$$

Theorem:

$\mathcal{H}_{p,m}$ is universal.

Pf:

- Let $k, \ell \in Z_P$ be two distinct keys.
- Given $h_{a,b}$, let

$$\begin{aligned} r &= (ak + b) \text{ mod } p, \\ s &= (a\ell + b) \text{ mod } p. \end{aligned}$$

- Then $r - s \equiv a(k - \ell) \text{ (mod } p))$

- For any $h_{a,b} \in \mathcal{H}_{p,m}$, distinct inputs k and ℓ map to distinct r and s modulo p .
- Each possible $p(p-1)$ choices for the pair (a,b) with $a \neq 0$ yields a different resulting pair (r,s) with $r \neq s$, since we can solve for a and b given r and s :

$$a = ((r-s)((k-\ell)^{-1} \bmod p)) \bmod p$$

$$b = (r-ak) \bmod p$$

- There are $p(p-1)$ possible pairs (r,s) with $r \neq s$, there is a 1-1 correspondence between pairs (a,b) with $a \neq 0$ and (r,s) , $r \neq s$.
- For any given pair of inputs k and ℓ , if we pick (a,b) uniformly at random from $Z_p^* \times Z_p$, the resulting pair (r, s) is equally likely to be any pair of distinct values modulo p .

- $\Pr[k \text{ and } \ell \text{ collide}] = \Pr_{r,s}[r \equiv s \pmod{m}]$
- Given r the number of s , such that $s \neq r$ and $s \equiv r \pmod{m}$, is at most

$$\begin{aligned}
 \lceil p/m \rceil - 1 &\leq ((p+m-1)/m) - 1 \\
 &= (p-1)/m \\
 \therefore s, s+m, s+2m, \dots, &\leq p
 \end{aligned}$$

- Thus,
- $$\Pr_{r,s} [r \equiv s \bmod m] \leq ((p-1)/m)/(p-1) = 1/m$$
- Therefore, for any pair of distinct key $k, \ell \in Z_p$:
- $$\Pr [h_{a,b}(k) = h_{a,b}(\ell)] \leq 1/m,$$

so that $\mathcal{H}_{p,m}$ is universal.

Hash with the multiply-shift method:

- Let $m = 2^\ell$ for some integer ℓ where $\ell \leq w$ and w is the number of bits in a machine word.
- Choose a fixed w -bit positive integer $a = A2^w$, where $0 < A < 1$ so that $0 < a < 2^w$.
- Assume that a key k fits into a single w -bit word and define:

$$h_a(k) = (ka \bmod 2^w) \gg (w - \ell)$$

$\mathcal{H} = \{h_a : a \text{ is odd, } 1 \leq a < m\}$, $2/m$ – universal

Hashing long inputs as vectors or strings :

- Can consider the class of vectors, such as vectors of 8-bit~ can be arbitrarily long.
- Number-theoretic approaches:
Extend the ideas used to design universal hash functions
- Cryptographic hash functions:
Use a hash function designed for cryptographic applications to design a good hash function for variable-length inputs.

Hashing long inputs as vectors or strings :

- Cryptographic hash functions take as input an arbitrary byte string and returns a fixed-length output.
- For example, the **NIST** standard deterministic cryptographic hash function **SHA-256** produces a **256-bit** (**32-byte**) output for any input.

$$h_a(k) = \text{SHA-256}(a || k) \bmod m$$

- Include instructions in CPU architectures to provide fast implementations of some cryptographic functions.
Advanced Encryption Standard (AES), the “**AES-NI**” instructions.

Open addressing:

- There is no list and no element stored outside the table.
- Advantage: avoid pointers, potentially yield fewer collisions and faster retrieval.

- $h: U \times \{0, 1, \dots, m-1\} \rightarrow \{0, 1, \dots, m-1\}$
- For every k , the probe sequence
$$\langle h(k, 0), h(k, 1), \dots, h(k, m-1) \rangle$$
is a permutation of $\langle 0, 1, \dots, m-1 \rangle$.
- Deletion from an open-address hash table is **difficult**.
- Thus chaining is more common when keys must be deleted.

HASH-INSERT(T, k)

```
1   $i = 0$ 
2  repeat
3       $q = h(k, i)$ 
4      if  $T[q] == \text{NIL}$ 
5           $T[q] = k$ 
6          return  $q$ 
7      else  $i = i + 1$ 
8  until  $i == m$ 
9  error “hash table overflow”
```

HASH-SEARCH(T, k)

```
1   $i = 0$ 
2  repeat
3       $q = h(k, i)$ 
4      if  $T[q] == k$ 
5          return  $q$ 
6       $i = i + 1$ 
7  until  $T[q] == \text{NIL}$  or  $i == m$ 
8  return  $\text{NIL}$ 
```

- Linear probing:
 - $h': U \rightarrow \{0, 1, \dots, m-1\}$ ~ an ordinary hash function (auxiliary hash function).
 - $h(k, i) = (h'(k) + i) \bmod m$

- Quadratic probing:
 - $h(k, i) = (h'(k) + c_1 i + c_2 i^2) \text{ mod } m$,
where h' is an auxiliary hash function, c_1 and $c_2 \neq 0$ and are constants.

- Double hashing:
 - $h(k, i) = (h_1(k) + ih_2(k)) \text{ mod } m$, where h_1 and h_2 are auxiliary hash functions.
 - $h_2(k)$ must be relative prime to m to search the whole hash table.
 - There are $O(m^2)$ probe sequences when m is prime or proper numbers, since each possible $(h_1(k), h_2(k))$ pair yields a distinct probe sequence.
 - Linear and Quadratic have $O(m)$ probe sequences.

0	
1	79
2	
3	
4	69
5	98
6	
7	72
8	
9	14
10	
11	50
12	

$$h(k, i) = (h_1(k) + ih_2(k)) \bmod m$$

Figure 11.5 Insertion by double hashing. Here we have a hash table of size 13 with $h_1(k) = k \bmod 13$ and $h_2(k) = 1 + (k \bmod 11)$. Since $14 \equiv 1 \pmod{13}$ and $14 \equiv 3 \pmod{11}$, we insert the key 14 into empty slot 9, after examining slots 1 and 5 and finding them to be occupied.

- Analysis of open-addressing hashing

$$\alpha = \frac{n}{m} : \text{load factor,}$$

with n elements and m slots.

- Thm:

Given an open-address hash table with load factor $\alpha = \frac{n}{m} < 1$, the expected number of probes in an **unsuccessful** search is at most $1/(1 - \alpha)$, assuming **independent uniform permutation hashing** and **no deletion**.

- Pf:
 - Define the r.v. X to be the number of probes made in an unsuccessful search.
 - Define A_i : the event there is an i -th probe and it is to an occupied slot.
 - Event $\{X \geq i\} = A_1 \cap A_2 \cap \cdots \cap A_{i-1}$

- $\Pr\{X \geq i\}$
- $= \Pr\{A_1 \cap A_2 \cap \cdots \cap A_{i-1}\}$
 $= \Pr\{A_1\} \cdot \Pr\{A_2 | A_1\} \cdot$
 $\Pr\{A_3 | A_1 \cap A_2\} \cdots \Pr\{A_{i-1} | A_1 \cap A_2 \cap \cdots \cap A_{i-2}\}$
- $\Pr\{A_1\} = n/m$
- The probability that there is a j th probe and it is to an occupied slot, given that the first $j-1$ probes were to occupied slots is $(n - j + 1)/(m - j + 1)$. Why?

- $\because n < m$, $(n - j)/(m - j) \leq n/m$ for all $0 \leq j < m$.

- $$\Pr[X \geq i] = \frac{n}{m} \cdot \frac{n-1}{m-1} \cdots \frac{n-i+2}{m-i+2}$$

$$\leq \left(\frac{n}{m}\right)^{i-1} = \alpha^{i-1}$$

$$E[X] = \sum_{i=1}^{\infty} \Pr[X \geq i] \leq \sum_{i=1}^{\infty} \alpha^{i-1} = \sum_{i=0}^{\infty} \alpha^i = \frac{1}{1-\alpha}$$



Cor: Inserting an element into an open-addressing hash table with load factor α requires at most $1/(1-\alpha)$ probes on average, assuming independent uniform permutation hashing and no deletion.

Theorem:

Given an open-address hash table with load factor $\alpha < 1$, the expected number of probes in a **successful** search is at most $\frac{1}{\alpha} \ln \frac{1}{1-\alpha}$, assuming **independent uniform permutation hashing with no deletion** and that each key in the table is equally likely to be searched for.

Pf: Suppose we search for a key k .

If k was the $(i+1)$ st key inserted into the hash table, the expected number of probes made in a search for k is at most


$$\frac{1}{1 - \frac{i}{m}} = \frac{m}{m-i}.$$

- Averaging over all n keys in the hash table gives us the average number of probes in a successful search:

$$\begin{aligned}
 & \frac{1}{n} \sum_{i=0}^{n-1} \frac{m}{m-i} = \frac{m}{n} \sum_{i=0}^{n-1} \frac{1}{m-i} = \frac{1}{\alpha} (H_m - H_{m-n}) \\
 & = \frac{1}{\alpha} \sum_{k=m-n+1}^m \frac{1}{k} \leq \frac{1}{\alpha} \int_{m-n}^m \frac{dx}{x} = \frac{1}{\alpha} \ln \frac{m}{m-n} = \frac{1}{\alpha} \ln \frac{1}{1-\alpha}
 \end{aligned}$$

11.5 Practical considerations

- Efficient hash table algorithms are not only of **theoretical interest**, but also of **immense practical importance**.
- Memory hierarchies: **cache memory, cache blocks**
- Advanced instruction sets



Linear probing with deletion

- $h(k, i) = (h_1(k) + i) \bmod m$
- The inverse function g maps a key k and a slot number q ($0 \leq q < m$) to the probe number that reaches slot q :

$$g(k, q) = (q - h_1(k)) \bmod m$$

- If $h(k, i) = q$, then $g(k, q) = i$, and so $h(k, g(k, q)) = q$.
- If $g(k', q) < g(k', q')$ then during the insertion of k' into the table, slot q was examined but found to be already occupied.

0	
1	
2	82
3	43
4	74
5	93
6	92
7	
8	18
9	38

(a)

0	
1	
2	82
3	93
4	74
5	92
6	
7	
8	18
9	38

(b)

- $$h(k, i) = h_1(k) + i,$$

$$h_1(k) = k \bmod 10$$
- (a) the hash table after inserting keys in the order 74, 43, 93, 18, 82, 38, 92
 - (b) the hash table after deleting the key 43 from slot 3

LINEAR-PROBING-HASH-DELETE(T, q)

```
1  while TRUE
2       $T[q] = \text{NIL}$                                 // make slot  $q$  empty
3       $q' = q$                                     // starting point for search
4      repeat
5           $q' = (q' + 1) \bmod m$                 // next slot number with linear probing
6           $k' = T[q']$                             // next key to try to move
7          if  $k' == \text{NIL}$ 
8              return                                // return when an empty slot is found
9          until  $g(k', q) < g(k', q')$         // was empty slot  $q$  probed before  $q'$ ?
10          $T[q] = k'$                             // move  $k'$  into slot  $q$ 
11          $q = q'$                                 // free up slot  $q'$ 
```

Theorem: If h_1 is 5-independent and $\alpha \leq 2/3$, then it takes **expected constant time** to search for, insert, or delete a key in a hash table using linear probing.

(M. Thorup. Linear probing with 5-independent hashing. <http://arxiv.org/abs/1509.04549>, 2015.)

Hash functions for hierarchical memory models: *wee hash function*

- $\text{swap}(x) = (x \gg (\frac{w}{2})) + (x \ll (\frac{w}{2}))$
 w : even word size (e.g., $w = 64$)
- $f_a(k) = \text{swap}((2k^2 + ak) \bmod 2^w)$, a is odd
- $f_a^{(r)}(k)$: result of iterating f_a a total of r times starting with t -bit input value k . 1-1 function?
- *wee hash function*:

$$h_{a,b,t,r}(k) = \left(f_{a+2t}^{(r)}(k + b) \right) \bmod m$$


WEE hash function for variable-length inputs

- Some k may be $> w$ -bit word in length or have variable length.
- $\text{chop}(k) = \langle k_1, \dots, k_u \rangle$,
 $u = \left\lceil \frac{t}{w} \right\rceil$,
padding 0's if necessary
- $h_{a,b,t,r}(k) =$
 $WEE(k, a, b, t, r, m)$

$$f_a(k) = \text{swap}((2k^2 + ak) \bmod 2^w)$$

$\text{WEE}(k, a, b, t, r, m)$

- 1 $u = \lceil t/w \rceil$
- 2 $\langle k_1, k_2, \dots, k_u \rangle = \text{chop}(k)$
- 3 $q = b$
- 4 **for** $i = 1$ **to** u
- 5 $q = f_{a+2t}^{(r)}(k_i + q)$
- 6 **return** $q \bmod m$



Perfect Hashing:

O(1) memory accesses are required for a search in the worst case

Figure 11.6 Using perfect hashing to store the set $K = \{10, 22, 37, 40, 60, 70, 75\}$. The outer hash function is $h(k) = ((ak + b) \bmod p) \bmod m$, where $a = 3$, $b = 42$, $p = 101$, and $m = 9$. For example, $h(75) = 2$, so key 75 hashes to slot 2 of table T . A secondary hash table S_j stores all keys hashing to slot j . The size of hash table S_j is m_j , and the associated hash function is $h_j(k) = ((a_j k + b_j) \bmod p) \bmod m_j$. Since $h_2(75) = 1$, key 75 is stored in slot 1 of secondary hash table S_2 . There are no collisions in any of the secondary hash tables, and so searching takes constant time in the worst case.

- **Perfect hashing :**
good for when the keys are static;
i.e. , once stored, the keys never
change, e.g. CD-ROM, the set of
reserved word in programming
language.

- **Thm** : If we store n keys in a hash table of size $m = n^2$ using a hash function h randomly chosen from a universal class of hash functions, then the probability of there being any collisions $< \frac{1}{2}$.
- **Proof:** Let h be chosen from an universal family. Then each pair collides with probability $1/m$, and there are $\binom{n}{2}$ pairs of keys.

Let X be a r.v. that counts the number of collisions. When

$$m = n^2, \quad E[X] = \binom{n}{2} \frac{1}{m} = \frac{n^2 - n}{2} \cdot \frac{1}{n^2} < \frac{1}{2}$$

By Markov's inequality, $\Pr[X \geq t] \leq E[X]/t$, and take $t = 1$.

Theorem:

If we store n keys in a hash table of size $m = n$ using a hash function h randomly chosen from universal class of hash functions, then $E[\sum_{j=0..m-1} n_j^2] < 2n$, where n_j is the number of keys hashing to slot j .

- Pf:
 - It is clear for any nonnegative integer a ,

$$a^2 = a + 2 \binom{a}{2}$$

$$\begin{aligned} & - \\ E\left[\sum_{j=0}^{m-1} n_j^2\right] &= E\left[\sum_{j=0}^{m-1}\left(n_j+2 \binom{n_j}{2}\right)\right] \\ &= E\left[\sum_{j=0}^{m-1} n_j\right]+2 E\left[\sum_{j=0}^{m-1} \binom{n_j}{2}\right] \end{aligned}$$

$$= E[n] + 2E\left[\sum_{j=0}^{m-1} \binom{n_j}{2}\right] = n + 2E\left[\underbrace{\sum_{j=0}^{m-1} \binom{n_j}{2}}_{\text{total number of collisions}}\right]$$

$$\binom{n}{2} \frac{1}{m} = \frac{n(n-1)}{2m} = \frac{n-1}{2}, \text{ since } m = n.$$

$$E\left[\sum_{j=0}^{m-1} n_j^2\right] \leq n + 2 \cdot \frac{n-1}{2} = 2n - 1 < 2n.$$



Corollary:

If store n keys in a hash table of size $m = n$ using a hash function h randomly chosen from a universal class of hash functions and we set the size of each secondary hash table to $m_j = n_j^2$ for $j = 0, \dots, m - 1$, then the expected amount of storage required for all secondary hash tables in a perfect hashing scheme is $< 2n$.

Corollary: Same as the above,

$$\Pr\{\text{total storage} \geq 4n\} < 1/2$$

Proof:

- By Markov's inequality, $\Pr\{X \geq t\} \leq E[X]/t$.

Take $X = \sum_{j=0}^{m-1} m_j$ and $t = 4n$:

$$\Pr\left\{\sum_{j=0}^{m-1} m_j \geq 4n\right\} \leq \frac{E\left[\sum_{j=0}^{m-1} m_j\right]}{4n} < \frac{2n}{4n} = \frac{1}{2}.$$

